

# Introduction to Apache Spark

Architecture, RDD, DataFrames, SQL & MLlib

Prasanth Kothuri, CERN

IT-Database Group

# Outline

What is Apache Spark?

Spark Abstractions

Spark Architecture

Spark Data APIs

- RDD

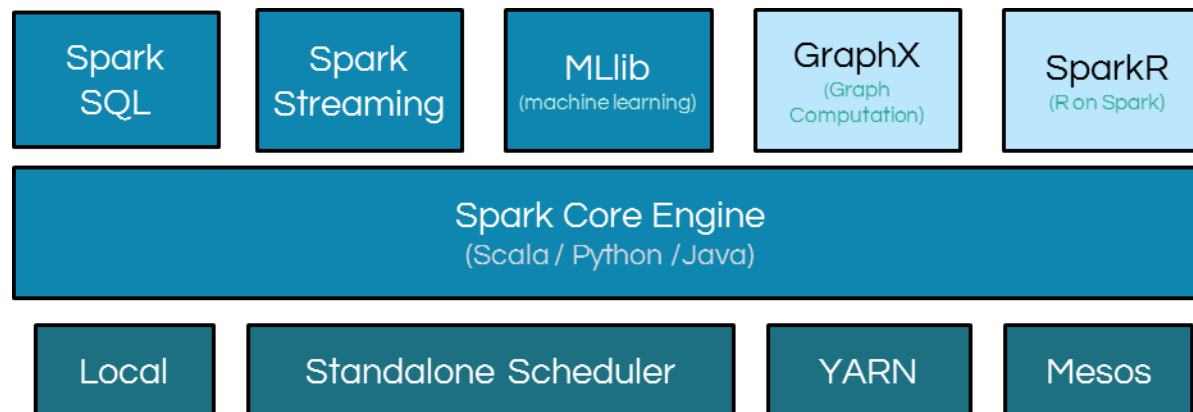
- Dataframe

- SQL

Spark Mlib

# What is Apache Spark?

- Apache Spark is an open-source parallel processing framework with expressive development APIs (in multiple languages) that allows for sophisticated analytics, real-time streaming and machine learning on large datasets
- Spark ecosystem



# Spark Abstractions

Two main abstractions of Spark

- RDD – Resilient Distributed Dataset
- DAG – Direct Acyclic Graph

# RDD

Resilient Distributed Datasets (RDDs) are the primary abstraction in Spark – a immutable distributed collection of records that can be operated on in parallel


There are currently two ways to create them:

- **parallelized collections** – take an existing python/scala collection and run functions on it in parallel
- **Hadoop datasets** – run functions on each record of a file(s) in Hadoop distributed file system

# RDD: Example

```
rdd = sc.parallelize(range(1,25),5)
```

rdd is split into  
partitions



**Partitioned:** RDD is partitioned and distributed across worker nodes of the cluster

**In-Memory :** RDD is stored in memory as much (size) and long (time) as possible


**Immutable:** does not change once created, can only be transformed into new RDDs

**Typed:** RDDs have types

**Cacheable:** hold all the data in a persistent storage like memory (preferable) or disk

# RDD: Actions and Transformations

- Two types of operations on RDDs:  
transformations and actions
- transformations – lazy operations that return another RDD
- actions – operations that trigger computation and return value



# DAG

**Direct Acyclic Graph** – sequence of computations performed on data

- Node – RDD partition
- Edge – transformation on top of data
- Acyclic – graph cannot return to the older partition
- Direct – transformation is an action that transitions data partition state (from X to Y)



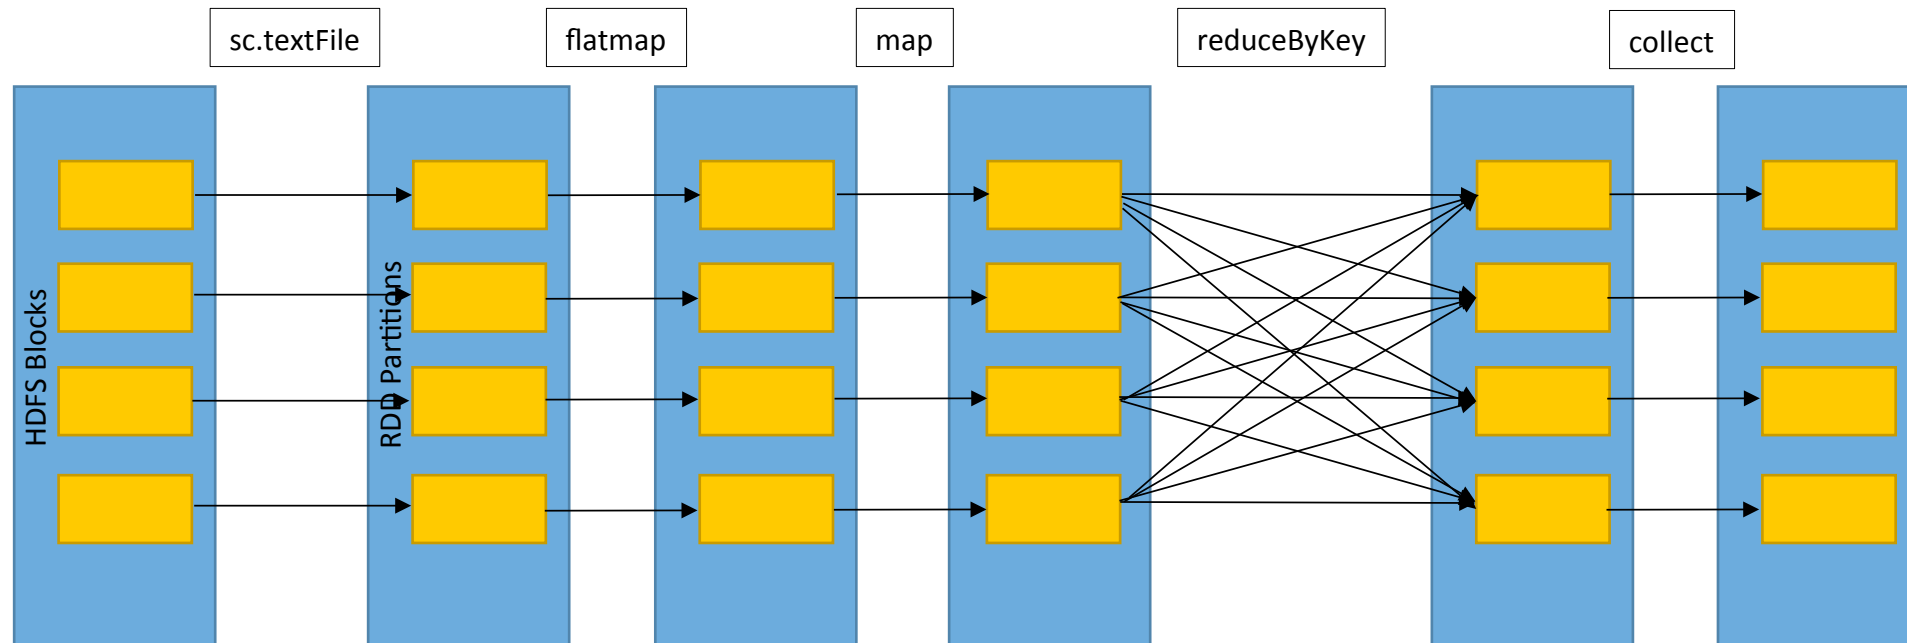
# DAG

## WordCount example

```
text_file = sc.textFile("hdfs:///user/pkothuri/sparkTraining/
LICENSE.txt")
counts = text_file.flatMap(lambda line: line.split(" ")) \
                    .map(lambda word: (word, 1)) \
                    .reduceByKey(lambda a, b: a + b)
for x in counts.collect():
    print x
```

# DAG

## WordCount example



# Spark WEB UI

Every SparkContext launches a web UI, by default on port 4040, that displays useful information about the application. This includes:

- A list of scheduler stages and tasks
- A summary of RDD sizes and memory usage
- Environmental information.
- Information about the running executors




## Stages for All Jobs

Completed Stages: 2

### Completed Stages (2)

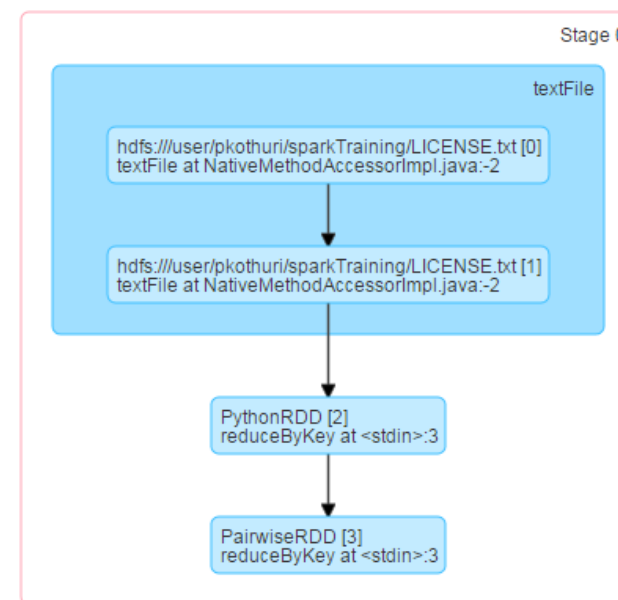
Stage Id	Description
1	collect at <stdin>:1
0	reduceByKey at <stdin>:3



## Details for Stage 0 (Attempt 0)

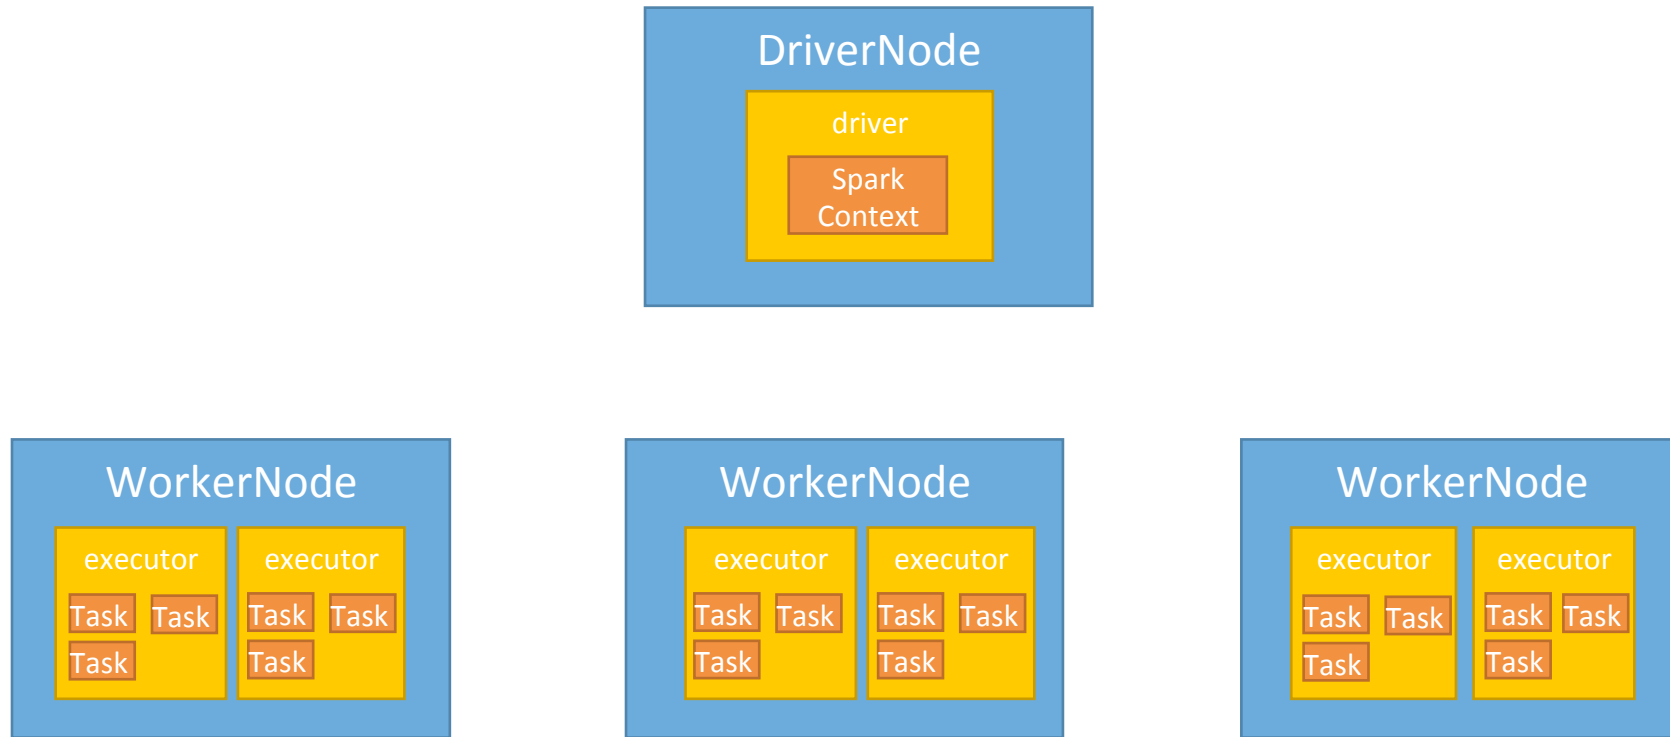
Total Time Across All Tasks: 7 s  
Locality Level Summary: Rack local: 2  
Input Size / Records: 5.5 KB / 202  
Shuffle Write: 9.8 KB / 24

[DAG Visualization](#)



# Spark Architecture

## Spark Cluster



# Spark Architecture

## **Driver**

- Entry point for Spark Shell (Scala, Python)
- SparkContext is created here and resides here
- Graph is built and submitted to DAGScheduler
- DAGScheduler divides it into stages and tasks
- Schedules tasks and controls their execution

# Spark Architecture

## **Executor**

- Reads data from HDFS (or external sources)
- Stores the data in cache in JVM heap or on HDDs
- Performs all data processing
- Writes data to HDFS (or external sources)

# Spark Architecture

## **Application Decomposition**

- Application
  - Single instance of SparkContext that stores some data processing logic and can schedule series of jobs, sequentially or in parallel
- Job
  - Complete set of transformations on RDD that finishes with action or data saving, triggered by the driver application

# Spark Architecture

## Application Decomposition

- Stage
  - Set of transformations that can be pipelined and executed by a single independent worker. Usually it is app the transformations between «read», «shuffle», «action» and «save»
- Task
  - Execution of the stage on a single data partition. Basic unit of scheduling

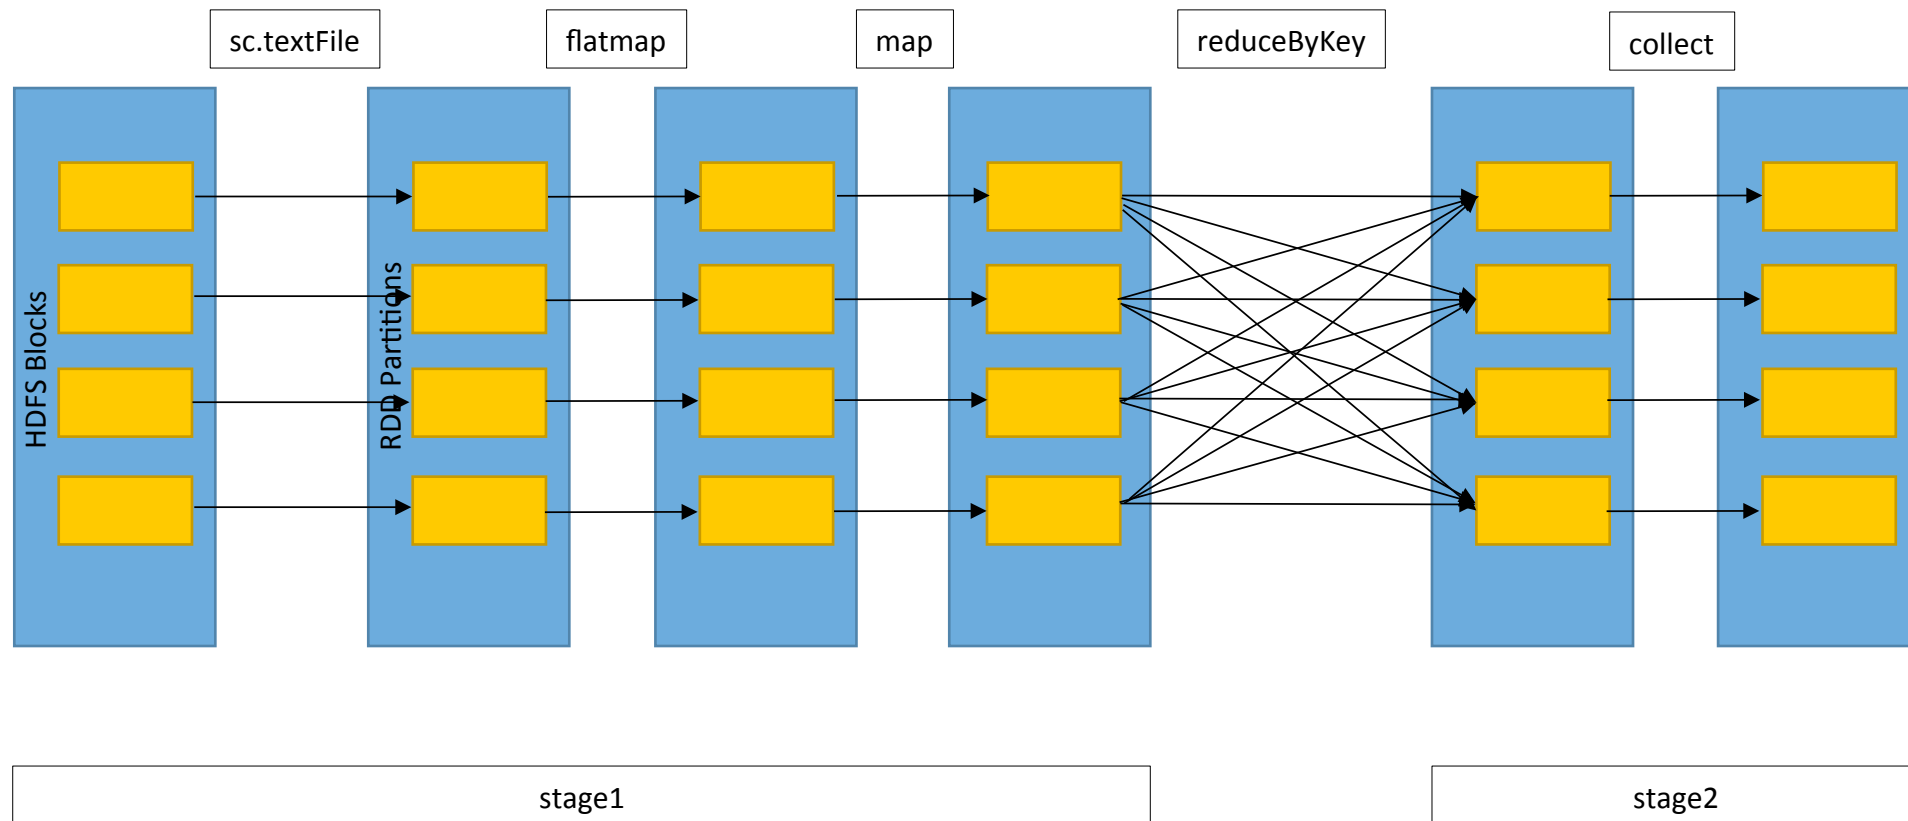


# DAG scheduler

- When an action is called on the RDD, Spark creates DAG and submits to the DAG scheduler
- The DAG scheduler divides operators into stages of tasks
- The stages are created based on the transformations, the narrow transformations are grouped (pipelined) into a single stage
- The DAG scheduler submits the stages to the task scheduler
  - The number of tasks depend on number of partitions
  - The number of tasks submitted depends on number of available executors

# DAG

## WordCount example




# SCHEDULING PROCESS

## RDD Objects

## DAG Scheduler


## Task Scheduler

## Executor




`Rdd1.join(rdd2)`  
`.groupBy(...)`  
`.filter(...)`

DAG




- Split graph into *stages* of tasks
- Stage boundaries are at wide transformations (shuffle)
- Submit each stage as ready

TaskSet



- Launches individual tasks
- Retry failed or straggling tasks

Task



- Execute tasks
- Store and serve blocks

- Build operator DAG



# Essential Spark Operations

## TRANSFORMATIONS

### General

- map
- filter
- flatMap
- mapPartitions
- groupBy
- sortBy
- flatMapValues
- groupByKey
- reduceByKey
- foldByKey
- sortByKey
- combineByKey

### Math / Stats

- sample
- sampleByKey
- randomSplit

### Set Theory

- union
- intersection
- subtract
- distinct
- cartesian
- zip
- join
- fullOuterJoin
- leftOuterJoin
- rightOuterJoin

### Data Structure / I/O

- keyBy
- zipWithIndex
- zipWithUniqueId
- zipPartitions
- coalesce
- repartition
- repartitionAndSortWithinPartitions
- pipe
- partitionBy

## ACTIONS

- reduce
- collect
- aggregate
- fold
- first
- take
- foreach
- top
- treeAggregate
- treeReduce
- foreachPartition
- collectAsMap
- keys
- values

- count
- takeSample
- max
- min
- sum
- histogram
- mean
- variance
- stdev
- sampleVariance
- countApprox
- countApproxDistinct

- takeOrdered

- saveAsTextFile
- saveAsSequenceFile
- saveAsObjectFile
- saveAsHadoopDataset
- saveAsHadoopFile
- saveAsNewAPIHadoopDataset
- saveAsNewAPIHadoopFile

# Transformations: map and filter

## MAP


Return a new RDD by applying a function to each element of this RDD

```
x = sc.parallelize([1,2,3,4,5])
y = x.map(lambda z: z * 2)
print(x.collect())
print(y.collect())
```

## FILTER

Return a new RDD that only has elements that pass the filter() function

```
x = sc.parallelize([1,2,3,4,5])
z = x.filter(lambda z: z % 2 != 0)
print(x.collect())
print(z.collect())
```



# Transformations: map and flatmap

## MAP

Return a new RDD by applying a function to each element of this RDD

```
x = sc.parallelize([3,4,5])
y = x.map(lambda z: [z,z*2])
print(x.collect())
print(y.collect())
```



## FLATMAP

Return a new RDD by applying a function to each element of the RDD, and then flattening the results. Also, function in flatMap can return a list of elements (0 or more)

```
x = sc.parallelize([3,4,5])
z = x.flatMap(lambda z: [z,z*2])
print(x.collect())
print(z.collect())
```



# Transformations: `reduceByKey` and `groupByKey`

## REDUCEBYKEY

Return a new RDD by combining the values with the same key with a given function

`[('a',3), ('b',4), ('c',5), ('a',4), ('b',-6)]`

reduceRDD

`[('a', 7), ('c', 5), ('b', -2)]`

```
x = sc.parallelize([('a',3), ('b',4), ('c',5), ('a',4), ('b',-6)])
y = x.reduceByKey(add)
print(x.collect())
print(y.collect())
```

## GROUPBYKEY

Return a new RDD by grouping the values with the same key

sourceRDD

groupRDD

`[('a', [3, 4]), ('c', [5]), ('b', [-6, 4])]`

```
x = sc.parallelize([('a',3), ('b',4), ('c',5), ('a',4), ('b',-6)])
I = x.groupByKey()
y = I.map(lambda z: (z[0],sum(z[1])))
print(x.collect())
print(y.collect())
```

# Actions

Actions trigger the computations and resulting result must fit in the driver JVM

<code>collect():</code>	returns all the elements of the RDD as an array to the driver, should be done after a filter or other operation
<code>count()</code>	returns the number of elements in the RDD
<code>countByKey()</code>	for pair RDDs returns (K, Int) pairs with the count of each key
<code>first()</code>	returns the first element of the RDD
<code>take(n)</code>	returns an array with first n elements
<code>saveAsTextFile()</code>	writes the elements of the RDD as a text file to HDFS or local filesystem
<code>getNumPartitions()</code>	returns the number of partitions of the RDD



# RDD: Demo

Login to lxplus

```
cd /eos/user/p/pkothuri
```

```
git clone https://github.com/prasanthkothuri/sparkTraining.git
```

Login to SWAN

<https://swan.cern.ch/>

open sparkTraining->notebooks->Tutorial\_RDD\_Final.ipynb

# Lifecycle of a Spark program

Create some input RDDs from external data or parallelize a collection in your driver program

Lazily transform them to define new RDDs using transformations like `filter()` or `map()`

Ask Spark to `cache()` any intermediate RDDs that will need to be reused

Launch actions such as `count()` and `collect()` to kick off a parallel computation, which is then optimized and executed by Spark


# Persistence and Cache

`persist()` is an action which triggers computation and persists a dataset in memory across operations

You can persist with different storage level options (e.g `MEMORY_ONLY` or `MEMORY_AND_DISK`)

`persist()` or `cache()` is key for iterative algorithms

`cache()` is same as `persist()` with `MEMORY_ONLY` storage option



# DataFrames

- An extension to the existing RDD API
- DataFrame is an RDD with schema
- DataFrames have numerous optimizations that make them much faster than RDD (predicate pushdown, bloom-filter)
- Write less code – solve problems concisely using dataframe functions
- Inspired by data frames in Python (pandas) and R

# Write Less Code: DataFrame vs RDD

## Using RDDs

```
w_rdd.map(lambda record: (record.NAMELAST, 1) ) \  
    .reduceByKey(add) \  
    .map(lambda (x,y):(y,x)) \  
    .sortByKey(False) \  
    .collect()
```

## Using DataFrames

```
w_df.groupBy("NAMELAST") \  
    .count() \  
    .orderBy("count") \  
    .show()
```

# Construct a DataFrame

```
# Create a DataFrame from json file
df = sqlContext.read.json("/tmp/read_my_shiny.json")
```

```
# Create a DataFrame by loading a parquet file
df = sqlContext.read.parquet("/tmp/path_to_the_parquet_file")
```

```
df = sqlContext.read.format('jdbc')
\ .options(driver='oracle.jdbc.driver.OracleDriver',url='jdbc:oracle:thin:username/
password@host:port:service_name',dbtable='table_name') \
.load()
```

```
# Convert a DataFrame
df = rdd.toDF()
```

# Schema Inference

DataFrames have schemas and can infer the schema from the type of the data being read

A Parquet file has a schema (column names and types) that DataFrames can use.

```
>>> df.printSchema()
root
|-- LocID: string (nullable = true)
|-- Location: string (nullable = true)
|-- VarID: string (nullable = true)
|-- Variant: string (nullable = true)
|-- Time: string (nullable = true)
|-- MidPeriod: string (nullable = true)
|-- SexID: string (nullable = true)
|-- Sex: string (nullable = true)
|-- AgeGrp: string (nullable = true)
|-- AgeGrpStart: string (nullable = true)
|-- AgeGrpSpan: string (nullable = true)
|-- Value: string (nullable = true)
```

What if the data doesn't have schema (e.g csv)

create an RDD of particular type using python namedtuple, dict and convert RDD to DataFrame

You can also specify the column names in .toDF() function

# DataFrame: Transformations and Actions

DataFrames are *lazy*. *Transformations* contribute to the query plan, but they don't execute anything.

*Actions* cause the execution of the query.

Execution of the query means:

- Spark initiates a distributed read of the data source
- The data flows through the transformations
- The result of the action is pulled back into the driver JVM.

## Transformation examples

- filter
- select
- drop
- intersect
- join

## Action examples

- count
- collect
- show
- head
- take



# Transformations: select(), filter() and show()

## filter()

The filter() method allows you to filter out rows from your results

## select()

similar to SQL SELECT, allows you to limit the results to specific columns

## show()

displays the first n elements in the DataFrame (n defaults to 20)

```
>>>w_df.filter(w_df.NAMELAST ==  
'KOTHURI').select(w_df.NAMELAST,w_df.NAMEFIRST,w_df.APPT_START_DATE).show()
```

```
+-----+-----+-----+  
|NAMELAST|      NAMEFIRST|APPT_START_DATE|  
+-----+-----+-----+  
| KOTHURI|      KRISHNA| 9/27/14 11:30|  
| KOTHURI|      SARYU| 9/27/14 11:30|  
| KOTHURI|VENKATAGOPALARAO| 9/27/14 11:30|  
+-----+-----+-----+
```

# Transformations: orderBy(), groupBy()

## orderBy()

The orderBy() method allows you sort the results

## groupBy()

groupBy() groups the elements by a specific column value, often used with count

```
>>>w_df.groupBy(w_df.NAMELAST) \
      .count() \
      .orderBy("count",ascending=False) \
      .show(10)
```

```
+-----+-----+
|NAMELAST|count|
+-----+-----+
|   Smith|25908|
| Johnson|21491|
|Williams|18228|
|   Brown|16804|
|   Jones|16023|
|   SMITH|14565|
|   Miller|12942|
|   Davis|12263|
| JOHNSON|12157|
|     Lee|10151|
+-----+-----+
```

only showing top 10 rows

# Transformations: Joins

r\_DF – is a DataFrame holding movie ratings [userId,movieId,rating,timestamp]

m\_DF – is a DataFrame holding movie information [movieId,title,genres]

These DataFrames can be joined as below to obtain number of reviews per genre

```
>>>r_DF.join(m_DF,r_DF.movieId == m_DF.movieId) \  
      .groupBy(m_DF.genres) \  
      .count() \  
      .orderBy("count",ascending =False) \  
      .show(5)
```

```
+-----+-----+  
|          genres|count|  
+-----+-----+  
|          Drama| 5832| | |
|          Comedy| 5648|  
|    Comedy|Romance| 3194|  
|    Drama|Romance| 2649|  
|Comedy|Drama|Romance| 2486|  
+-----+-----+  
only showing top 5 rows
```

# Spark SQL and DataFrames

DataFrames and Spark SQL are essentially tied to each other

- The DataFrames API provides a programmatic interface for interacting with data
- Spark SQL provides a SQL-like interface
- Whatever you can do in DataFrames, you can do in Spark SQL and vice versa

# Spark SQL contd.

Spark SQL allows you to manipulate distributed data with SQL queries. Currently, two SQL dialects are supported.

- If you're using a Spark SQLContext, the only supported dialect is "sql", a rich subset of SQL 92.
- If you're using a HiveContext, the default dialect is "hiveql", corresponding to Hive's SQL dialect. "sql" is also available, but "hiveql" is a richer dialect.

# Spark SQL contd.

- You issue SQL queries through a `SQLContext` or `HiveContext`, using the `sql()` method.
- The `sql()` method returns a `DataFrame`.
- You can mix `DataFrame` methods and SQL queries in the same code.
- To use SQL, you must either:
  - query a persisted Impala or Hive table, or
  - make a table alias for a `DataFrame`, using `registerTempTable()`

# Spark SQL - Example

To issue SQL against an existing DataFrame, create a temporary table, which essentially gives the DataFrame a *name* that's usable within a query.

```
>>> df = sqlContext.read.parquet("/tmp/WH_VRecord/part-r-00000-5396c70a-ff5b-4dda-9306-3a5e8bd9167a.gz.parquet")

>>> df.registerTempTable("VistorRecords")

>>> sql_df = sqlContext.sql("SELECT NAMELAST,NAMEFIRST,APPT_START_DATE,APPT_END_DATE FROM VistorRecords")

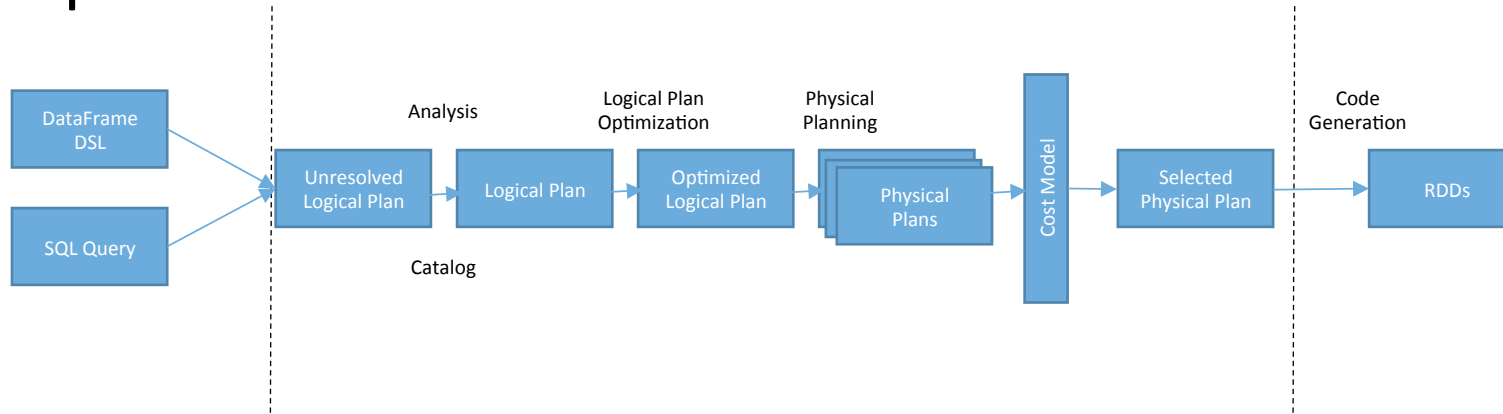
>>> sql_df.show(5)
```

NAMELAST	NAMEFIRST	APPT_START_DATE	APPT_END_DATE
Adamopoulos	Stella	5/1/15	5/1/15
Brosman	Muriel	5/1/15	5/1/15
Brumfield	Avery	5/1/15	5/1/15
Chipman	Catherine	5/1/15	5/1/15
Chubb	Steven	5/1/15	5/1/15

```
only showing top 5 rows
```

# Catalyst : Spark's Optimizer

- Spark SQL uses catalyst to optimize all the queries written both in spark sql and dataframe dsl



## Analysis

- phase where attribute references or relations are resolved
- e.g: column validity, column type
- catalog object tracks the tables in all data sources



# Catalyst : Spark's Optimizer

## Logical Optimizations

- Standard rule-based optimizations
- e.g: predicate pushdown, project pruning, null propagation etc

## Physical Planning

- generated one or more physical plans
- cost model is used to select a plan

## Code Generation

- generate java bytecode to speed up execution

## Further Reading

- <https://databricks.com/blog/2015/04/13/deep-dive-into-spark-sqls-catalyst-optimizer.html>

# Dataframe and SPARK SQL: Demo

Login to lxplus

```
cd /eos/user/p/pkothuri
```

```
git clone https://github.com/prasanthkothuri/sparkTraining.git
```

Login to SWAN

<https://swan.cern.ch/>

open sparkTraining->notebooks->Tutorial\_DataFrame\_Final.ipynb

# Spark MLlib

## Why Apache Spark for Machine Learning?

- Bigger than memory datasets
  - Able to train a model on large scale dataset
- General Purpose
  - Apart from the libraries for commonly used algorithms, libs for advanced data preparation, feature engineering etc
- Compatibility
  - Support for multiple languages and integrate well with python libs like pandas, scikit-learn etc

# Dataframe and RDD based API

## DataFrame-based API (spark.ml)

- Easier to construct a machine learning pipeline
- Flexible and versatile API compared to spark.mllib
- Will reach feature parity with spark.mllib in the next releases

## RDD-based (spark.mllib)

- Original Machine Learning API
- No new features, only bugfixes
- Will be removed in Spark 3.0

# Spark MLlib – Main Concepts

## DataFrame

- Same as the Dataframe from Spark SQL / Dataframe API
- Used to hold ML dataset

## Transformer

- Transforms one DataFrame into another DataFrame
- Examples
  - Feature transformer appends new column (features) to DataFrame
  - Learning model transforms a DataFrame with features into a DataFrame with predictions

## Estimator

- Abstracts the learning algorithm; accepts a DataFrame and produces a Model

## Pipeline

- A sequence of transformers and estimators to process and learn from data

## Parameter

- Specifying parameters for transformers and estimators

# Spark MLlib utilities and algorithms

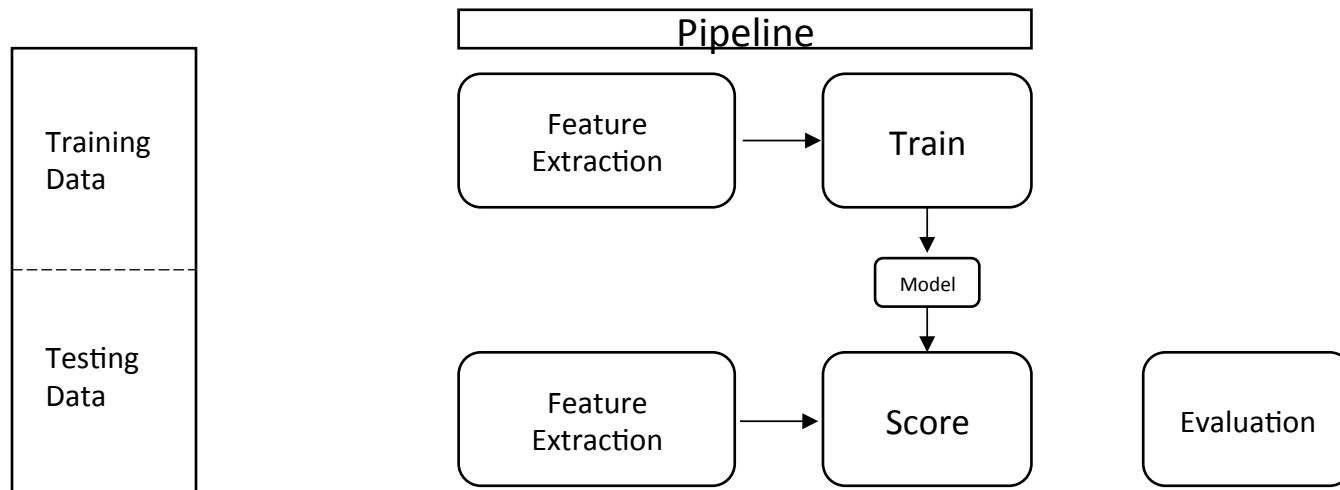
## Distributed pre-processing workflow utilities

- Feature Engineering
  - Extraction:- Extracting features from raw data
    - Word2vec
    - countvectorizer
  - Transformations:- modifying, converting or scaling features
    - Tokenizer
    - StringIndexer
    - Standardization
    - Normalization
    - VectorAssembler
  - Selectors:- selecting a subset from a larger set of features

## High performant ML algorithms (classification, regression, clustering etc)

- Full list - <http://spark.apache.org/mllib/>

# Model Lifecycle



# SPARK MLlib: Demo

Login to lxplus

```
cd /eos/user/p/pkothuri
```

```
git clone https://github.com/prasanthkothuri/sparkTraining.git
```

Login to SWAN

<https://swan.cern.ch/>

open sparkTraining->notebooks->Tutorial-ML-Final.ipynb



# Conclusion

- We have covered the spark concepts; abstraction and architecture
- Introduced to Spark data APIs – RDD, DataFrame and Spark SQL
- Demonstration of using Spark data APIs for exploratory data analysis and data analytics
- Introduced to Spark Mllib for scalable machine learning
- Understood how spark can aid in distributed computing of VERY large datasets
- Several ways to interact with spark – spark-shell, pyspark, spark-submit and jupyter notebooks

# Further study

- Slides and Notebooks for Hands On
  - <https://github.com/prasanthkothuri/sparkTraining>
- 2016 IT DB Hadoop tutorials
  - <https://github.com/prasanthkothuri/hadoop-tutorials-2016>
- Coursera course – Big Data Analysis with Scala and Spark
  - <https://www.coursera.org/learn/scala-spark-big-data>
- CERN Technical Training on Apache Spark in June